

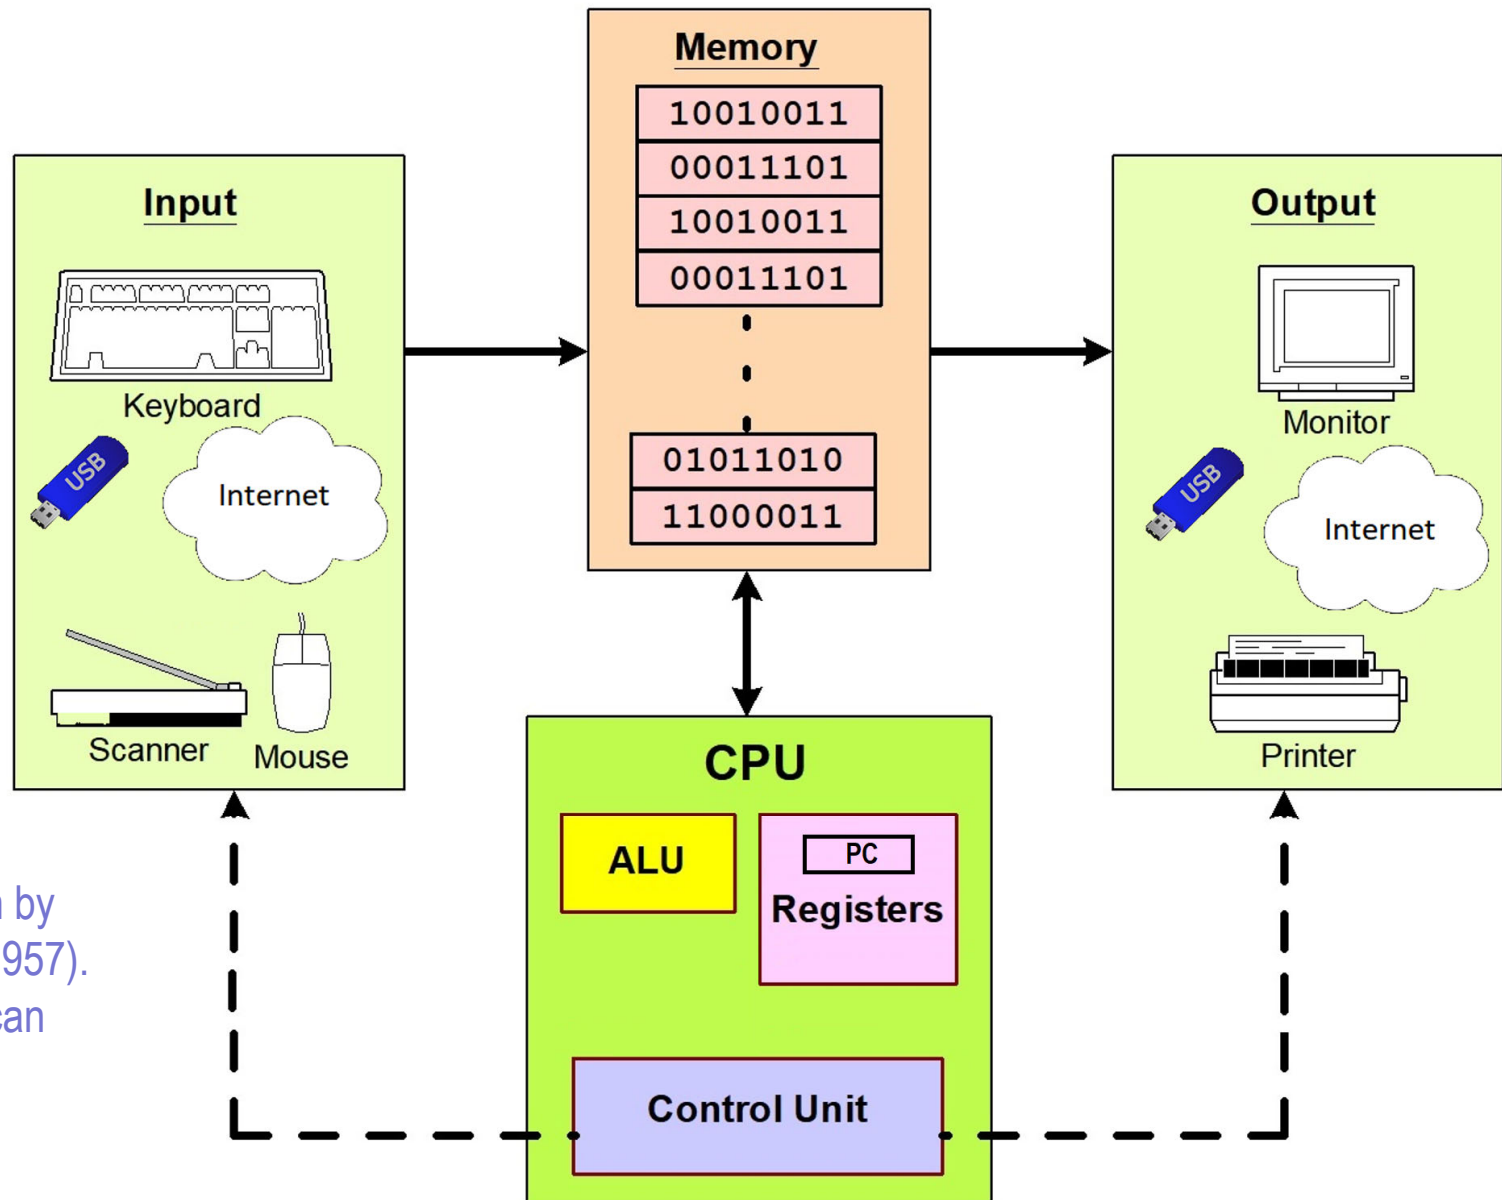
A decorative L-shaped line consisting of a vertical black line on the left and a horizontal grey line extending to the right, intersecting at a small black crosshair.

EECE7376: Operating Systems Interface and Implementation

Introduction



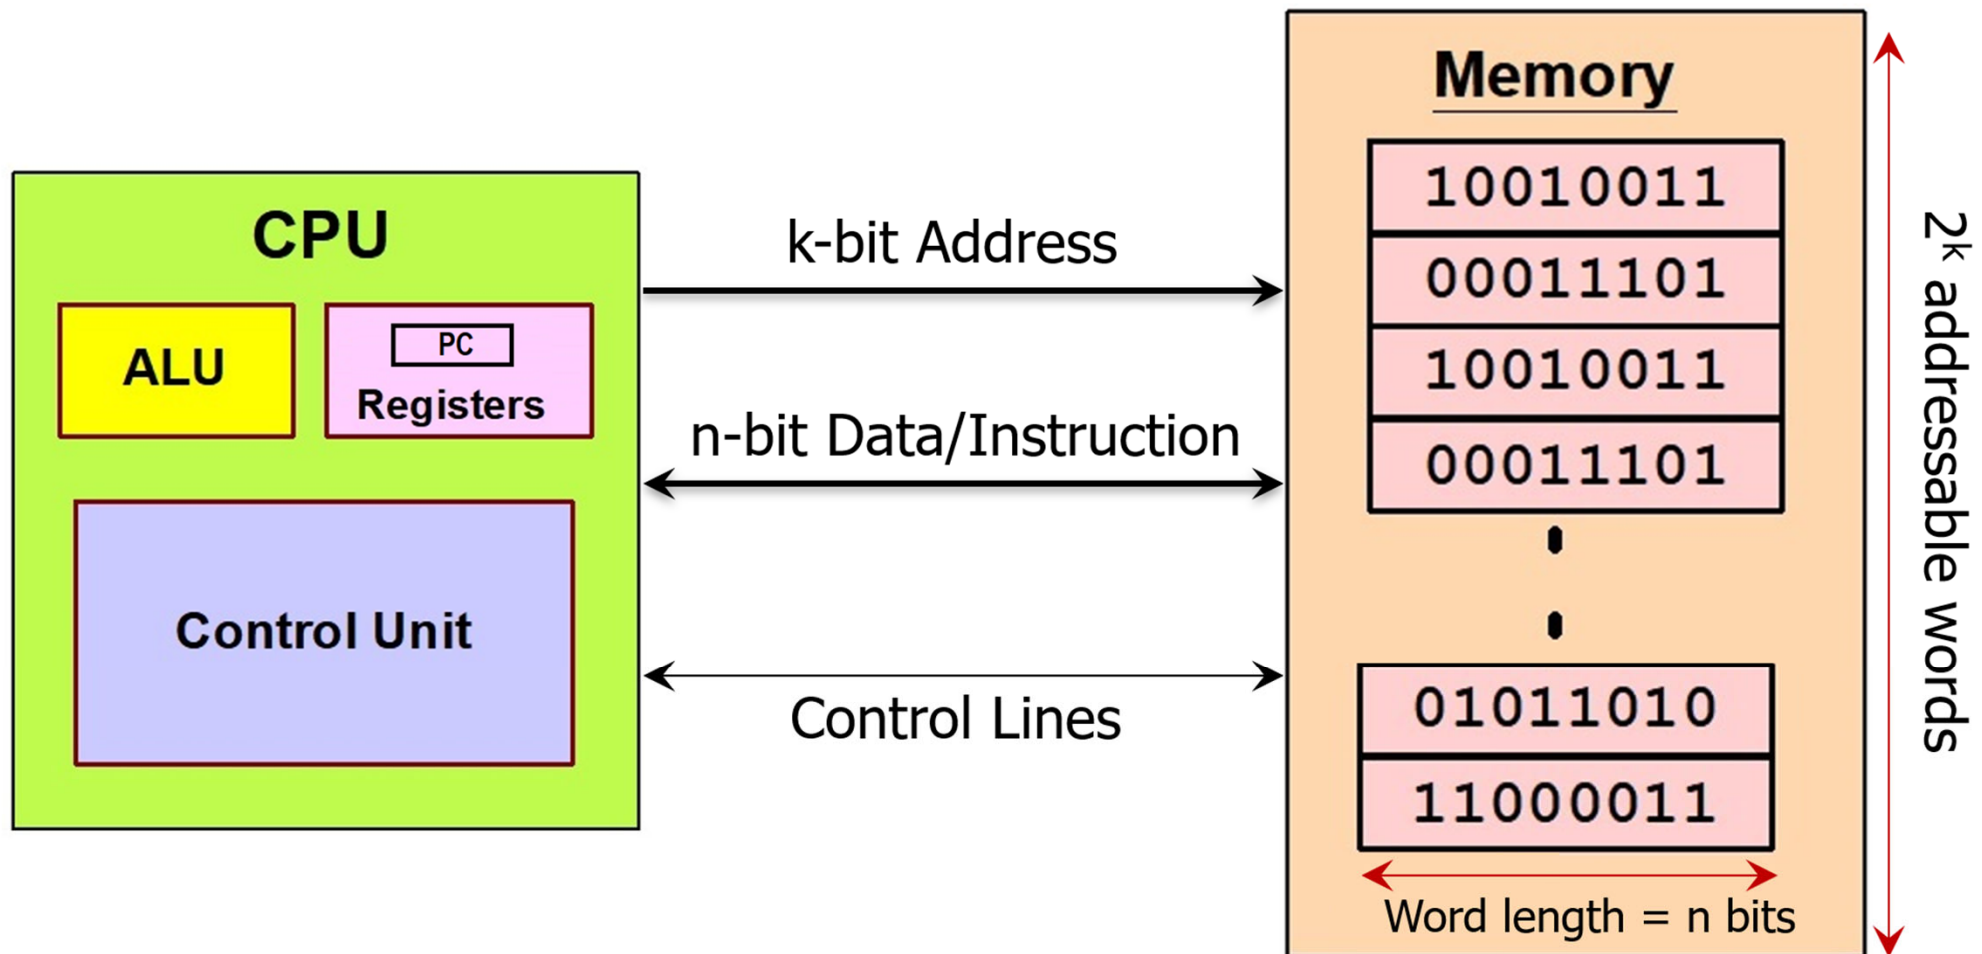
The Von Neumann Computer Model



It is a computer architecture based on a 1945 description by John Von Neumann (1903-1957). He was a Hungarian-American mathematician, physicist, inventor, computer scientist



Processor/Memory Interface




When a computer starts, what is the address of the first instruction to be executed and is this first instruction part of the operating system?



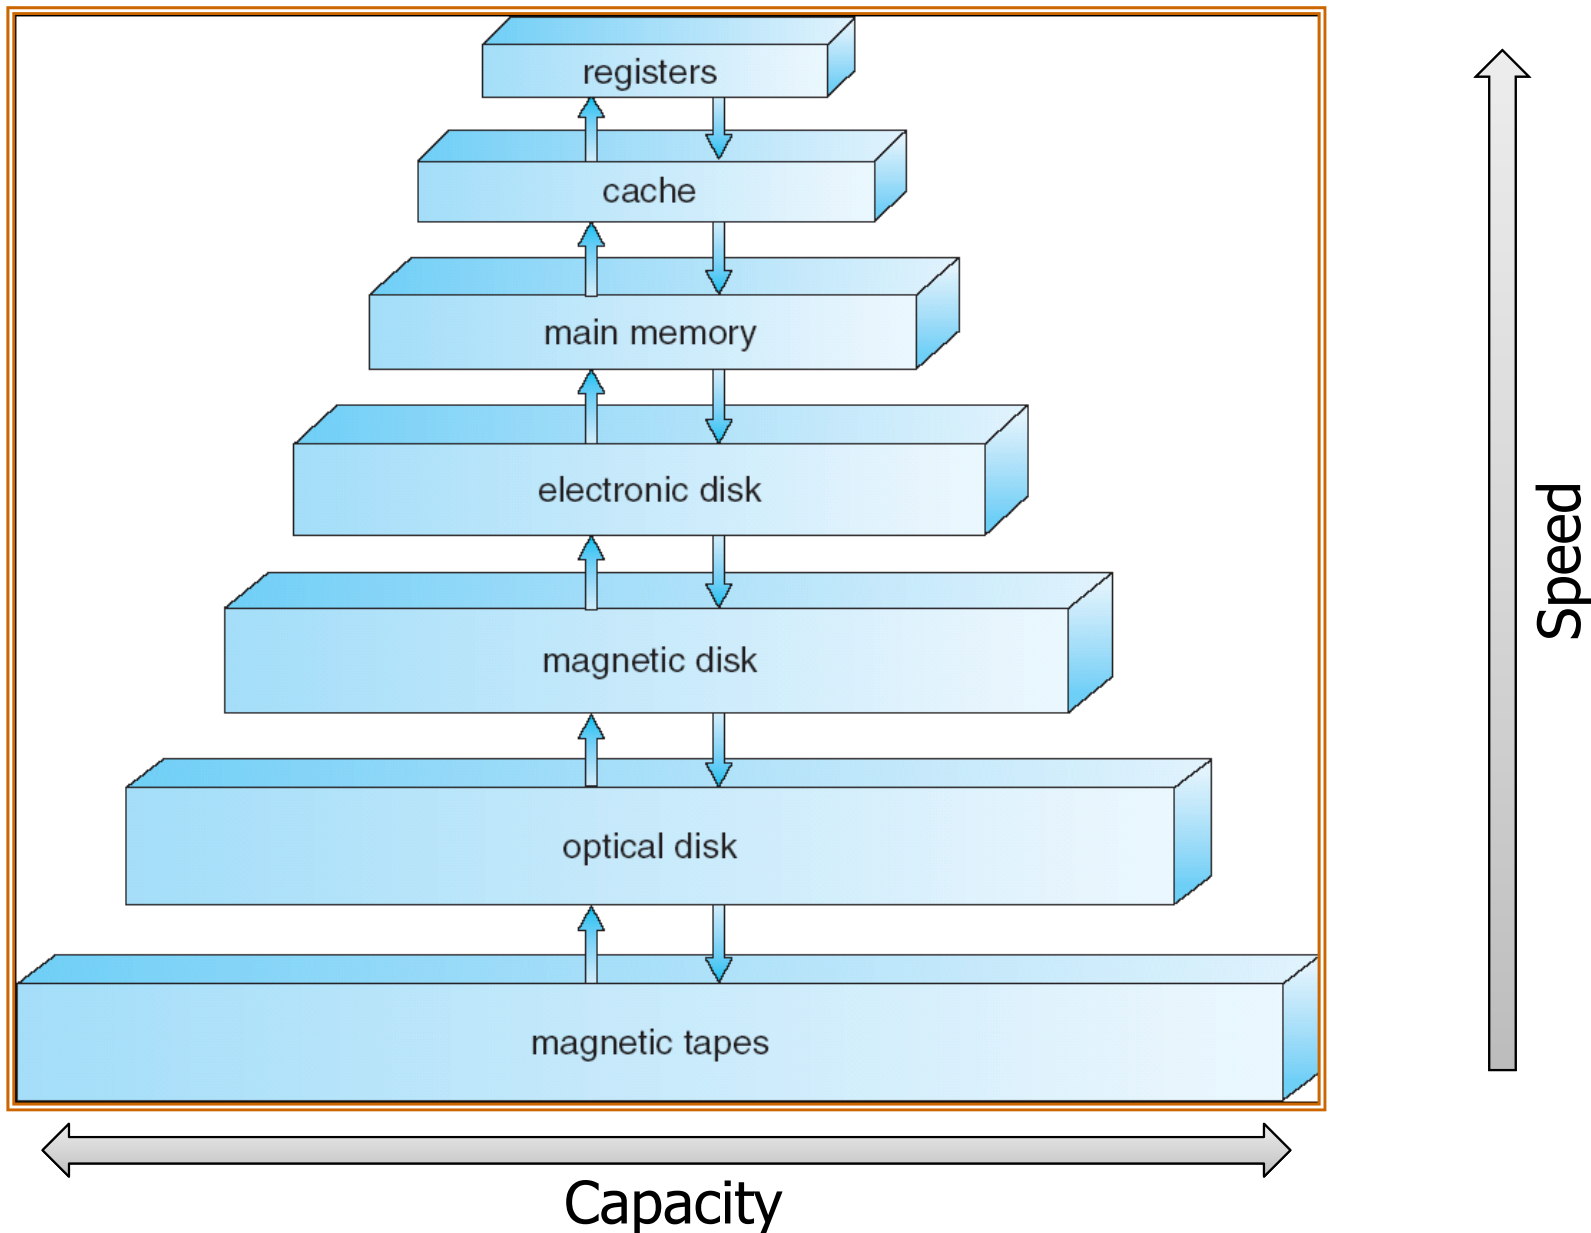
Memory Layout of a Program

- Text: program code
- Static data:
 - Global and static variables C++
- Stack:
 - Local variables of functions
- Dynamic data: heap
 - In C++, created by **new** and removed by **delete**





Memory/Storage Hierarchy





An Instruction Execution Steps

1. Fetch and PC update

- Use the program counter (PC) to supply the instruction address and fetch instruction from memory.
- Update the PC for the next fetch (e.g., $PC \leftarrow PC + 4$)

2. Decoding decodes instruction and reads operands

- Extract opcode: determine what operation should be done
- Extract operands: register numbers or immediate from the instruction

3. Execution

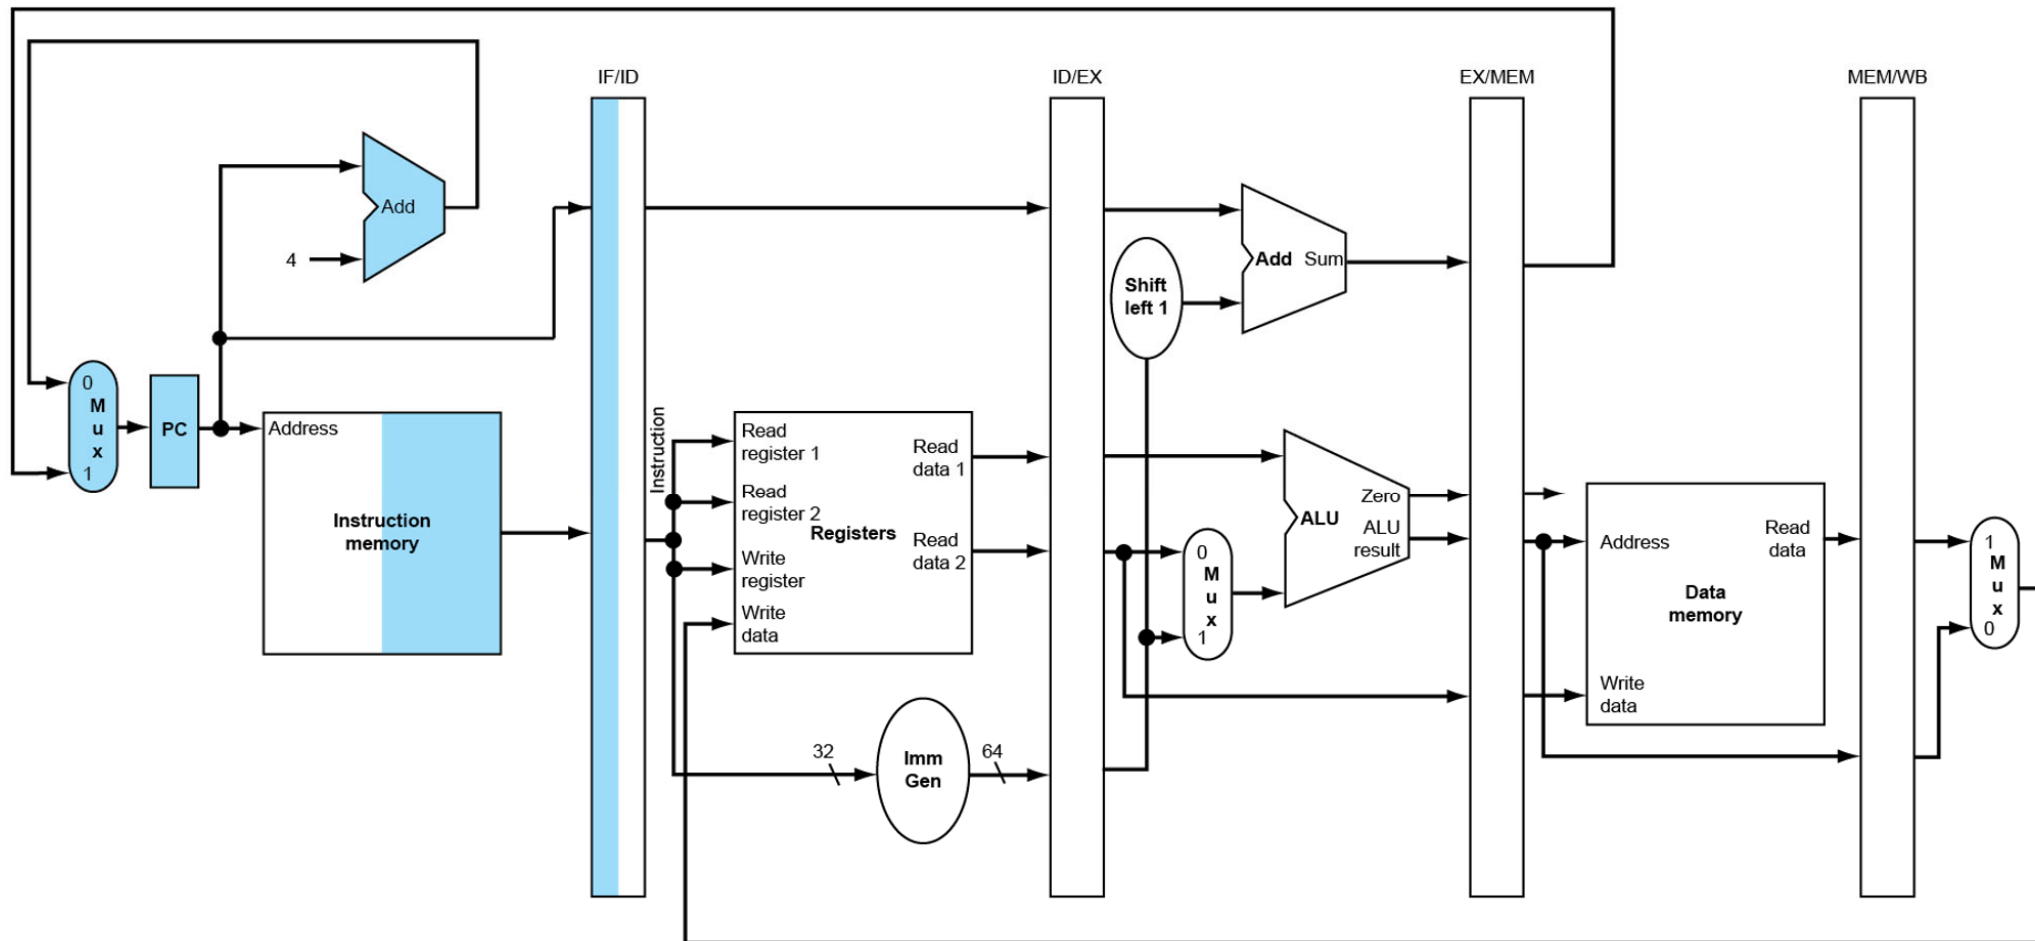
- Use ALU to calculate one of the following tasks:
 - Arithmetic or logical operations on registers contents.
 - Calculating the memory address for load/store
 - Calculating the branch target address
- Access memory for load/store

4. Write back one of the following results

- $PC \leftarrow$ target address for jump and taken branch
- Update destination register for R-type, addi, or lw instructions.
- Update destination memory for the sw instruction.



RISC Pipelined Microarchitecture



A processor guarantees **atomic execution** of each single instruction (either executes it as a whole or none of it).



Sample Assembly Program

- The following RISC-V program prints the contents of an array of one-byte integers on the screen.

```
.data
    v: .byte 2, 7, 1, 4, 3
.text
main:
    li t0, 0      # Array index
cond:
    li t1, 5      # Array size
    bge t0, t1, endloop
body:
    la t2, v      # Array base address
    add t2, t2, t0 # base + offset
    lb t2, 0(t2)
```

```
li a7, 1 # print int
mv a0, t2
ecall
```

```
addi t0, t0, 1
j cond
```

```
endloop:
    li a7, 10 # exit
    ecall
```

What if the exit ecall is missing?

Why does the program need the OS to handle this output operation instead of writing directly to the screen?



C Pointers Review

- A pointer is merely an address of where a datum or structure is stored
 - Pointers operations are based on the type of entity that they point to.
 - To declare a pointer, use `*` preceding the variable name: `int *p;`
- To set a pointer to a variable's address use `&` before the variable as in `p = &x;`
 - `&` means "return the memory address of"
 - in this example, `p` will now point to `x`
- If you access `p`, you merely get the address
- To get the value that `p` points to, use `*` as in `*p`
 - `*p = *p + 1;` will add 1 to `x`
- `*` is known as the **indirection** (or **dereferencing**) operator because it requires a **second access**.
- `*p` and `x` are known as aliases, two ways of accessing the same memory location.



Pointers Example 1

```
int main() {  
    int x = 1, y = 2;  
    int z[3] = {10, 20, 30};  
    int *p;  
    p = &x;           // p now points at the location where x is stored  
    y = *p;           // set y equal to the value pointed to by p, or y = x  
    printf("x= %d   y= %d\n", x, y); // x=1 y=1  
    *p = 0;           // now change the value that p points to to 0, so now x = 0  
                        // but will that change y's value?  
    printf("x= %d   y= %d\n", x, y); // x=0 y=1  
    p = &z[0];        // now p points at the first location in the array z, z[0]  
    *p = *p + 1;       // the value that p points to (z[0]) is incremented  
    ++p;              // now p points at the second location in the array z, z[1]  
    *p = *p + 1;       // the value that p points to (z[1]) is incremented  
    printf("z[0]= %d   z[1]= %d\n", z[0], z[1]); // z[0]=11 z[1]=21  
    return 0;  
}
```

By how much is the content of **p** incremented
after the **p++** operation?



Arrays and Pointers

Example: `int z[3] = {10, 20, 30};`

- In C, the name of the array is actually a pointer to the first array element (i.e., `z = &z[0]`)
- Therefore, there are two ways to access the first element of array `z`: either `z[0]` or `*z`
- What about accessing `z[1]`?
 - We can do `z[1]` as usual, or we can add 1 to the location pointed to by `z`, that is `*(z+1)`. **Is this equivalent to `*(++z)`?**
- While we can update a pointer value (like `++p` in Example 1), we cannot update the value of `z` (e.g., `++z`), otherwise we would lose access to the first array location since `z` is our array variable.
 - Notice in Example 1, if an `int` size is 4 bytes then when we did `p++`, `p`'s value is incremented by 4. This is because `p` is of type "`int *`".
- If `p` is pointing to doubles, the increment `++p` would increment `p` by 8 bytes away, and by 1 if it points to `char` type.



Iterating Through the Array

- The following are two ways to print an array's contents:

```
int j;  
int z[3] = {10, 20, 30};  
for(j = 0; j < 3; j++)  
    printf("%d \n", z[j]);
```

```
int *p;  
int z[3] = {10, 20, 30};  
for(p = z; p < z + 3; p++)  
    printf("%d \n", *p);
```

- For the code on the right:
 - p started by having the value of z , that is the address of $z[0]$.
 - The loop iterates while $p < z + 3$; where $z+2$ is the address of the last element of the 3-element array z .
 - $p++$ increments the pointer to point at the next element in the array.
- The code on the right is more efficient as the only operation inside the loop is incrementing p . While the code on the left increments j and then add $j*4$ to the base address of the array.
- *Note:* $++(*p)$; increments the value of the element to which p points. While $*(++p)$; increments the pointer to point at the next array element and then return the value of that element.



Pointers Example 2

```

int main()
{
    int x[4] = {12, 20, 39, 43}, *y;
    y = x;                                // y points to the beginning of the array
    printf("%d \n", x[0]) ;                // outputs 12
    printf("%d \n", *y) ;                  // also outputs 12
    printf("%d \n", *y+1) ;                 // outputs 13 (12 + 1)
    printf("%d \n", *(y+1));               // outputs x[1] or 20
    y+=2;                                  // y now points to x[2]
    printf("%d \n", *y) ;                  // prints out 39
    *y = 38;                               // changes x[2] to 38
    printf("%d \n", *y-1) ;                 // prints out x[2] - 1 or 37
    y++;                                  // sets y to point at the next array element
    printf("%d \n", *y) ;                  // outputs x[3] (43)
    (*y)++;                               // sets what y points to, to be 1 greater
    printf("%d \n", *y) ;                  // outputs the new value of x[3] (44)
    return 0;
}

```



Passing Arrays to Functions

- When an array is passed to a function, what is being passed is a pointer to the array
 - In the formal parameter list, you can either specify the parameter as an array or a pointer

```
int array[100];
```

```
...
```

```
afunction(array);
```

```
...
```

```
void afunction(int *a) {...}
```

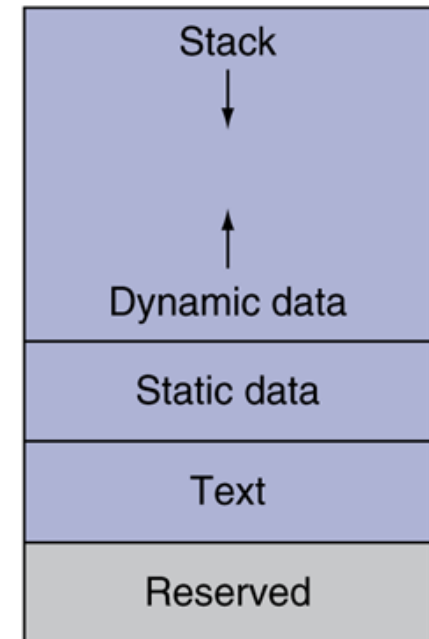
or

```
void afunction(int a[]) {...}
```



Dynamic Memory Allocation (1 of 2)

- To this point, we have been declaring pointers and having them point at already created variables.
- An important use of pointers is to create dynamic structures.
 - structures that can have data added to them or deleted from them such that the amount of memory being used is equal to the number of elements in the structure.
 - this is unlike an array which is static in size.
- Creating and maintaining dynamic data structures requires dynamic memory allocation – the ability for a program to obtain more memory space at execution time and to release the space that is no longer needed.





Dynamic Memory Allocation (2 of 2)

- In C, the **malloc** and **free** *library functions* are essential to dynamic memory allocation.
- **malloc** is used to request memory space enough to hold a specific data type or an array of the data type.
- Dynamically allocate memory to an integer

```
int *ptr = (int *) malloc(sizeof(int));
```

- Dereference the pointer to access the memory

```
*ptr = 8;           // assigns 8 to memory
```

- At the end free up the memory for reuse

```
free(ptr);          // returns memory to the system.
```



Dynamic Memory Operators for Arrays

- Dynamically allocate memory to an integer array with 10 elements

```
int *arr = (int *) malloc(10 * sizeof(int));
```

- Dereference the pointer to access the memory

```
arr[0] = 8; // access element at index 0
```

- At the end free up the memory for reuse

```
free(arr); // free up memory.
```



Pointers Example 3

//Program to read and calculate the average of class grades

```
float AverageGrades(int grades[], int n) {
    float sum = 0;
    for(int i = 0; i < n; i++) sum += grades[i];
    return sum/n;
}

int main() {
    int* gradesP;
    int ClassSize;
    printf("What is the class size? ");
    scanf("%d", &ClassSize);
    gradesP = (int *) malloc(ClassSize * sizeof(int));

    for(int i = 0; i < ClassSize; i++) {
        printf("What is grade # %d ? ", i);
        scanf("%d", &gradesP[i]);
    }
    printf("Grades average = %f \n", AverageGrades(gradesP, ClassSize));
    free(gradesP);
    return 0;
}
```



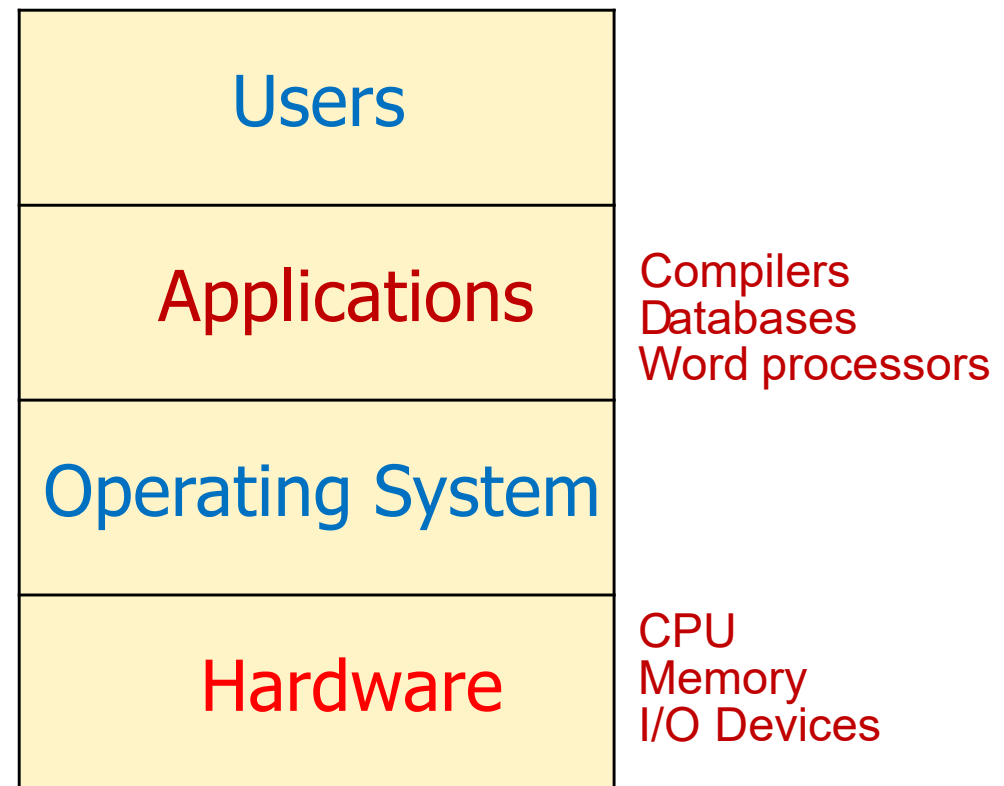
Linux and C Resources

- COE Linux gateway:
<https://wiki.coe.neu.edu/xwiki/bin/view/Main/#HLinuxHelp>
For accounts support: <https://coe.northeastern.edu/computer/>
To access from off campus:
<https://wiki.coe.neu.edu/xwiki/bin/view/Main/VPN/>
- Installing Windows Subsystem for Linux (WSL) and then the Ubuntu Linux distribution:
 - <https://docs.microsoft.com/en-us/windows/wsl/install>
 - <https://ubuntu.com/tutorials/ubuntu-on-windows#1-overview>
 - You will need to install the gcc compiler by running:
`sudo apt install g++ or sudo apt install gcc`
- Incomplete introduction to the C programming environment from the book Appendix F:
 - <https://pages.cs.wisc.edu/~remzi/OSTEP/lab-tutorial.pdf>



What is an Operating System?

- The OS is a software that acts as an intermediary between a user of a computer and the computer hardware.
- It is responsible for making sure the computer system operates correctly, efficiently, fairly, reliably, and in an easy-to-use manner.
- It carries its task by providing:
 - Virtualization
 - Resource management
 - Security and Protection





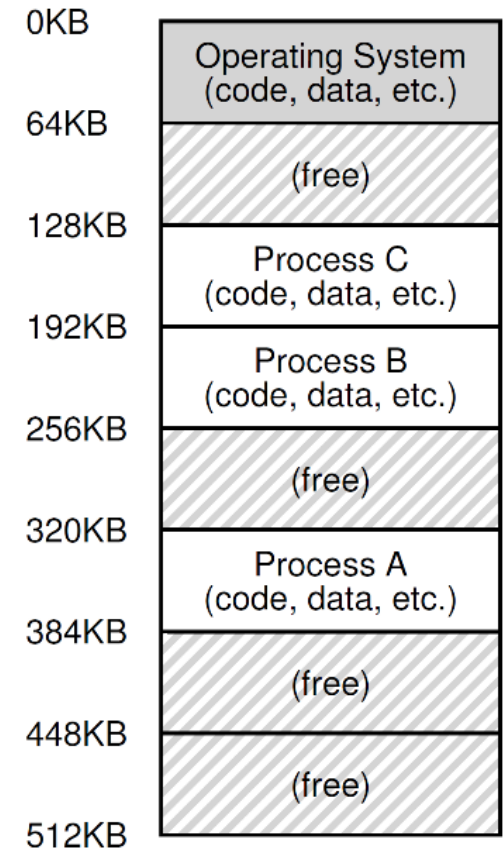
Virtualization

- That is, the OS takes a physical resource (such as the processor, or memory, or a disk) and transforms it into a more general and easy-to-use virtual form of itself.
- The OS manages and abstracts the low-level hardware, so that, for example, a word processor need not concern itself with which type of disk hardware is being used.



OS Kernel and System Calls

- The OS kernel is a special program that provides services to running programs (processes).
- When a process needs to invoke a kernel service, it invokes a procedure call in the operating system interface. Such a procedure is called a **system call**.
- The system call enters the kernel; the kernel performs the service and returns. Thus, a processor alternates between executing in **user space** and **kernel space**.
- System calls are usually called through the programming language standard library. The library makes different devices look the same by providing higher-level abstractions





Resources Management

- Virtualization allows many programs to concurrently:
 - run (thus sharing the CPU),
 - access their own instructions and data (thus sharing memory), and
 - access devices (thus sharing for example the disks),
- The OS is sometimes known as a *resource manager*.
- Advantages of managing computer resources:
 - **Protect** applications from one another
 - Provide **efficient** and **fair** access to resources



CPU Virtualization Experiment

- Code: github.com/remzi-arpacidusseau/ostep-code/tree/master/intro
- Run four background processes of the `cpu.c` code
`./cpu A & ./cpu B & ./cpu C & ./cpu D &`
- *Observation*: Even though we have only one processor, somehow all four of these processes seem to be running at the same time.
- *Conclusion*: The OS, with some help from the hardware, oversees giving the illusion that the system has a very large number of virtual CPUs. Allowing many programs to seemingly run at once is what we call virtualizing the CPU.
- *Question*: if two programs start to run at the same time, which should run first? The OS as a **resource manager** needs to answer this by providing a **policy**.
- *Linux tips*:
 - In WSL Ubuntu, to open a new terminal select Ubuntu from the down arrow on the top bar.
 - `ps -A` //lists all running processes in all terminals
 - `killall cpu` //kills the above running background cpu processes
 - `kill <process number>` //kill specific process using its number.



CPU Virtualization Code

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[])
{
    if (argc != 2){
        fprintf(stderr, "usage: cpu <string>\n");
        exit(1);    }

    char *str = argv[1];
    while (1)      {
        printf("%s\n", str);
        Spin(1);    }
    return 0;
}
```



Memory Virtualization Experiment

- Code: github.com/remzi-arpacidusseau/ostep-code/tree/master/intro
- Run two background processes of the `mem.c` code (you might need to add `(int *)` before `malloc`): `./mem 1 & ./mem 200 &`
- *Observation:*
 - Each of the above process allocates memory at the same address (e.g., `0x200000`)*, and yet each seems to be updating the value at `0x200000` independently! It is as if each running program has its own private memory, instead of sharing the same physical memory.
- *Conclusion:*
 - Each process accesses its own private **virtual address space**, which the OS and hardware somehow maps onto the physical memory of the machine.
 - A memory reference within one running program does not affect the address space of other processes.
 - The reality, however, is that physical memory is a shared resource, managed by the operating system (the **resource manager**).

*For this to happen, your OS needs to be set to disable address-space layout randomization(ASLR). In Ubuntu this can be done by: (replace 0 with 2 to enable it)

```
echo 0 | sudo tee /proc/sys/kernel/randomize_va_space
```



Memory Virtualization Code

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: mem <value>\n");
        exit(1);    }
    int *p;
    p = malloc(sizeof(int));    assert(p != NULL);
    printf("(%)d) addr pointed to by p: %p\n", (int) getpid(), p);
    *p = atoi(argv[1]); // assign value to addr stored in p
    while (1) {
        Spin(1);
        *p = *p + 1;
        printf("(%)d) value of p: %d\n", getpid(), *p);    }
    return 0; }
```



Concurrency Experiment

- Code: github.com/remzi-arpacidusseau/ostep-code/tree/master/intro
- Compile the `threads.c` code using:

```
g++ threads.c -pthread -o threads
```
- Run it once with argument 1,000 and several times with argument 100,000.
- *Observation:*
 - With 1000, the final value of the counter is 2000 as expected, as each thread incremented the counter 1000 times.
 - With 100000, every time we get a different wrong value for the counter (we might sometimes get the correct 200000 value).
- *Conclusion:*
 - Both threads **share** the same data segment.
 - The reason for these odd outcomes is that instructions are executed one at a time. Incrementing the counter takes three instructions: load the value of the counter from memory into a register, increment the register content, and finally store this content back into memory.
 - As these three instructions do not execute **atomically** (all at once), strange things happen. It is the problem from a **race condition** with **concurrency**.



Concurrency Code

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"
#include "common_threads.h"
volatile int counter = 0;
int loops;

void *worker(void *arg) {
    int i;
    for (i = 0; i < loops; i++)
        counter++;

    return NULL;
}
```

```
int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr,
            "usage: threads <loops>\n");
        exit(1); }
    loops = atoi(argv[1]);
    pthread_t p1, p2;
    printf("Initial value : %d\n", counter);
    Pthread_create(&p1, NULL, worker, NULL);
    Pthread_create(&p2, NULL, worker, NULL);
    Pthread_join(p1, NULL);
    Pthread_join(p2, NULL);
    printf("Final value   : %d\n", counter);
    return 0;
}
```




Persistence

- We need hardware and software to be able to store data **persistently**; vs. the **volatile** storage of the computer RAM.
 - Example saving files in hard disks or solid-state drives (SSDs).
- The software in the OS that usually manages these storage drives is called the **file system** that provides the user with several systems calls.
 - Unlike the abstractions provided by the OS for the CPU and memory, it does not create a private, virtualized storage drive for each application. Rather, it is assumed that **users want to share information** that is in files (among processes and sometimes among other users).



Persistence Experiment

- Code: github.com/remzi-arpacidusseau/ostep-code/tree/master/intro
- Run one process of the `io.c` code.
- *Observation:*
 - We used three system calls (open, write, and close) to create a file and write a text in it.
 - With the help of the OS **file system** and the various **device drivers**, the file and its content are accessible through directory and file names, regardless of the technology of the physical storage device.
- *Conclusion:*
 - The OS provides a standard and simple way to access devices through its **system calls**.



Persistence Code

```
#include <stdio.h>
#include <unistd.h>
#include <assert.h>
#include <fcntl.h>
#include <sys/stat.h>
#include <sys/types.h>
#include <string.h>

int main(int argc, char *argv[]) {
    int fd = open("h.txt",
                  O_WRONLY | O_CREAT | O_TRUNC, S_IRUSR | S_IWUSR);
    assert(fd >= 0);
    char buffer[20];
    sprintf(buffer, "hello world!!\n");
    int rc = write(fd, buffer, strlen(buffer));
    assert(rc == (strlen(buffer)));
    fsync(fd);
    close(fd);
    return 0;
}
```



OS Design Goals (1 of 2)

- Build up some **abstractions** in order to make the system convenient and easy to use.
- High **performance**; another way to say this is our goal is to minimize the overheads of the OS.
- **Security** against external malicious applications is critical, especially in these highly-networked times.
- Provide **protection** between processes. Protection include the mechanisms that allow the implementation of security.
 - The heart of one of the main principles underlying an operating system, which is that of **isolation**; isolating processes from one another is the key to protection.



OS Design Goals (2 of 2)

- The operating system must also run non-stop; when it fails, all applications running on the system fail as well. Therefore, operating systems often strive to provide a high degree of **reliability**.
- **Energy-efficiency** is important in our increasingly green world.
- **Mobility** is increasingly important as OSes are run on smaller and smaller devices.



Reading List

- From Arpaci-Dusseau textbook:
 - Section 2.6 Some History